

004 — Tier 0: SpaceX Anchor, SOTP across Space / Connectivity / AI

SPCX's live \$2.07 trillion capitalisation is 312.8× the annualised operating income of Connectivity — its only profitable segment. No conglomerate discount exists to measure, and no admissible comparator exists for any of the three regimes: the partition closes at zero of eleven.

Constitution pin	1.5.0
As of	2026-09-18
Universe	SPCX 100%
Generated	2026-09-19 14:00 UTC

1. The anchor's entire price rests on the one segment that earns
2. The finding survives every allocation the market could plausibly make
3. The three regimes fail comparability for three different reasons
4. Space is a captive supplier and carries no transaction price at all
5. Connectivity earns the profit, and its growth is a mix shift
6. The AI segment's own headline metric cannot be turned into a value
7. SPCX files a three-level revenue decomposition, and every level closes exactly
8. The capital goes to the segment that loses the most
9. The live price requires perpetual growth equal to the required return
10. The constitution's own anchor does not reconcile with the IPO price
11. Four bounds stand, and a fifth was retracted because the filing answers it
12. All six pillars hold, and three of them hold only as bounds
13. Twelve findings were run against this thesis, four of them high
14. What 005, 006 and 009 may price off, and what they may not

TIER 0 SPACEX ANCHOR · THESIS 004 · AS OF 2026-09-18

The anchor's entire price rests on the one segment that earns

\$2.0718T

LIVE MARKET CAP
13,567,041,680 shares ×
\$152.71, close 2026-09-18

312.8×

MARKET CAP ÷
CONNECTIVITY OP.
INCOME
\$2,071.8bn over \$6,624M
annualised from a filed quarter

0 of 11

ADMISSIBLE
COMPARATORS
the partition admits none, for
three different reasons

\$ (143)M

CONSOLIDATED OP.
LOSS, Q2 2026
segments sum exactly to (143);
the served layer returns +143

The news: SpaceX is the largest company ever to list, and it lists as a **broadband operator carrying two loss-making ventures**. Space runs a **\$(2,168)M** annualised operating loss and AI **\$(5,028)M**. Connectivity is the only segment with a positive operating result — **\$1,656M in the quarter** — and the whole **\$2.07 trillion** is **312.8×** that segment's annualised income.

Why that inverts the standard question. A sum-of-the-parts analysis normally asks whether the parts are worth more than the whole, and reports the gap as a conglomerate discount. **That question is not available here.** The parts do not sum to less than the whole under any admissible multiple — the whole already requires the one profitable segment to carry a ~313× operating multiple. **What this means for the position** is that the discount is not small, it is absent, and the live question becomes a different one: **is the market crediting the two loss-making segments with any value at all?**

The anchor's job is to bound that question rather than answer it. This thesis produces three valuation regimes — one per segment — each a **MODELED** range with its boundary named, and **no blended multiple and no point estimate**. It is the layer beneath the sector theses: 005, 006 and 009 each price a company that consumes what this anchor establishes.


And the anchor's most uncomfortable result is a refusal. The comparability partition admits **zero of eleven** named comparators — not a search that failed, but **three structurally different reasons** that one screen would not have found. A downstream thesis that borrows one of the eleven has re-opened a closed class without saying so, and the failure would be invisible because that thesis would look internally consistent.

How to read this report. The finding and why it survives every allocation come first, then the three boundaries, the segments one at a time, the capital structure and the market print — **including two corrections this thesis made to its own artifacts while producing it**. **The recommendation is not a position. This thesis produces no trades by design. It produces the floor every downstream thesis stands on.**

The finding survives every allocation the market could plausibly make

The 313× figure has an obvious objection: it attributes the **whole** market capitalisation to one segment. If the other two carry value, Connectivity's implied multiple falls. **This page runs that objection to its limit, and the finding holds across the entire plausible range.**

Value granted to Space + AI	Residual for Connectivity	Implied multiple	Market cap granted away
\$0	\$2,071.8bn	312.8×	0%
\$0.5T	\$1,571.8bn	237.3×	24%
\$1.0T	\$1,071.8bn	161.8×	48%
\$1.5T	\$571.8bn	86.3×	72%



Connectivity's implied operating multiple (x, dimensionless) as an increasing share of the \$2.0718 trillion market capitalisation is granted to the Space and AI segments. Even granting the two loss-making segments **\$1.5 trillion — 72% of the price — Connectivity still carries 86.3× its annualised operating income.** Derived from filed segment income at a 3M rate ×4, stated as an annualisation and not a forecast. Source: [SPCX 10-Q p.30](#).

Why the ladder is the evidence rather than the 313×. A single attribution invites the reply that the analyst chose it. **The ladder removes the choice:** at every allocation a reader might defend, Connectivity carries a multiple with no terrestrial analogue. **This means** the conclusion is not an artefact of one attribution — **the robustness across the range is the finding**, which is why the anchor reports a ladder instead of a number.

Three of the four standard SOTP components do not exist here in the form the template assumes. The **conglomerate discount** cannot be measured as a discount and is reported as its inverse. The **control and float discount** is made harder by the Cursor consideration being **Class A stock**, so the two cannot be separated from the transaction the anchor now includes. **Unallocated corporate cost is not separable** — SPCX files no such line, and the anchor reports it as a bound rather than allocating it arbitrarily.

The three regimes fail comparability for three different reasons

The partition is the deliverable. A boundary is a **partition** — a name is in or out, class named — not a graded score, because **a score lets each downstream thesis pick its own threshold. Admitted: zero of eleven.**

Segment	Excluded names	Classes	Why no multiple can be borrowed
Space	RKLB · FLY	loss_making, P11	No earnings — both run losses. RKLB carries two independent exclusions : it also acquired IRDM
Connectivity	IRDM · GSAT · ASTS · VSAT · SATS	P11, PARTIAL, DA-24	The price is a deal spread — IRDM \$54.00 , GSAT \$90.00 /share, where 64% of GSAT revenue is one customer who is also its acquirer
AI	MSFT · GOOG · NVDA · VRT	non_disclosure	No discrete segment exists — MSFT's and GOOG's compute are cost centres ; NVDA is the supplier, not the operator

The three reasons are different, and that difference is the argument. `loss_making` is an **earnings** problem — no numerator. `P11` is a **price-formation** problem — a numerator measuring deal terms, not a business. `non_disclosure` is a **segment-boundary** problem — the business exists, the disclosure does not. **This means** no single screen produces this partition: one for unprofitable peers misses IRDM, for clean prices misses MSFT, for segment disclosure misses RKLB.

The trap the live feed makes worse. IRDM and GSAT are now **priceable** — but **a live multiple on a merger spread measures deal terms, not the business. Verifiable is not admissible. What this means for downstream work:** the anchor is **004's own segment data or nothing** — and since SPCX is in neither 006's nor 009's universe, this is the programme's only valuation of SPCX's Connectivity and AI segments.

SPACE · THE SEGMENT THAT HAS NO PRICE

Space is a captive supplier and carries no transaction price at all

Space is simultaneously the highest-gross-margin and the worst-operating-margin business in the company, and the 122.14-point gap between those two facts is fully decomposed by two filed expense lines.

Space, Q2 2026	Value	Note
Revenue	\$962M	+29.0% year on year — customer launches only
Gross margin	65.80%	the highest in the company
Operating margin	-56.34%	the worst in the company
The swing	122.14 pp	R&D 111.85 pp + SG&A 10.29 pp — decomposed exactly
Ex-R&D operating result	+\$534M	633 gross profit – 99 SG&A — the launch business earns
Comparability	none	captive_integrated — no transaction price exists

The +\$534M is the segment's load-bearing number. It is what remains when Starship development is stripped out, and it establishes that **launch carries standalone value** rather than being a subsidised internal function. **The limit travels with it:** the filing does **not** split Starship from Falcon R&D, so the split is **MODELED** and the artifact says so.

Most of what Space does is never sold. Of **38 launches** in Q2 2026, **10 carried a customer payload — 26.3%**. The other **28 were internal Starlink deployments**, and for those the issuer states: *"the Company does not recognize any inter-segment revenue, rather those launch costs are capitalized in satellites in Property, plant, and equipment, net."* **This means** the segment's revenue line measures a **minority of its output**, and Space's contribution to the group is **capitalised into Connectivity's asset base** — which is why 001's own finding was that value migrated **out of launch**. Launch counts: [SPCX 10-Q p.35](#); the inter-segment rule: [p.36](#).

Why no external multiple may be borrowed, stated as a measurement problem rather than a coverage one. A competitor's launch margin is a **price against a cost**. Space's is a **residual customer book**. They are not two measurements of one quantity — **comparing them compares a price to a non-price**. The value therefore comes from segment **contribution**, and the DA-06 flag travels with the row into every downstream thesis that cites it.

CONNECTIVITY · THE ONLY SEGMENT THAT EARNS

Connectivity earns the profit, and its growth is a mix shift

Connectivity is the highest-margin operator in the universe and the only profitable SPCX segment. Its operating income grew **faster than its revenue**, which is the operating-leverage case demonstrated on the margin rather than argued.

Connectivity, Q2 2026	Value	Change
Revenue	\$4,291M	+65.8% — 54.9% of consolidated
Operating income	\$1,656M	+79.4%
Operating margin	+38.59%	the only positive segment
Rate spread	+13.6 pp	income outgrew revenue by 13.6 points
Enterprise & Government	\$1,806M	+108.3%
Consumer	\$2,485M	+44.4% — the managed channel grew 2.4× faster
Starlink subscribers	12.0M	from 6.0M — +100.0% point-to-point; the issuer cites 101.2% on an average-subscriber basis
Starlink ARPU	\$66/month	from \$85 — -22.4%

The mix is moving toward the higher-value channel, and that is the growth story. Enterprise & Government — aviation, maritime, land mobility, fixed sites and government entities — grew **2.4× faster than consumer**, both figures filed as dimensions of the same axis. **This means** the extension is **a mix shift rather than a rising tide**, and it is the more durable of the two: a managed contract is a relationship, not a consumer price point.


And the ARPU decline is real — but ARPU is not the segment's monetisation metric, and the two facts are different. The issuer states the cause directly: consumer revenue grew with *"101.2% growth in Starlink subscribers, offset by a 22.4% decline in Starlink subscriber ARPU, primarily due to international expansion and the addition of lower priced service plans."* That is a **genuine mix shift toward cheaper international plans**, and the report does not discount it. **What ARPU cannot measure is the rest of the segment:** its denominator *"does not include managed enterprise and government customers"*, so it **excludes the channel growing at +108.3%**. **This means** two separate findings, kept separate — the consumer price decline is real, **and** ARPU is structurally blind to the fastest-growing part of Connectivity. DA-10 is registered because a service-line ARPU is not a company ARPU. Cause: [SPCX 10-Q p.43](#); definition: [p.36](#).

AI · THE SEGMENT WHOSE HEADLINE METRIC CANNOT BE CONVERTED

The AI segment's own headline metric cannot be turned into a value

AI is the largest drag on the consolidated line and the fastest-growing revenue line at once, and its headline metric is a capacity figure that no disclosed price converts into money.

AI, Q2 2026	Value	Note
Revenue	\$2,561M	+247.5% — contaminated: xAI merged 2026-02-02, prior periods recast
Operating loss	\$(1,257)M	-49.08% — the largest single drag on consolidation
Headline metric	1.4 GW	nameplate compute draw, from 0.4 GW — installed capacity, not consumption
Advertising, separable	\$710M H1	21.0% of the segment — the one part of AI that decomposes
Comparability	none	MSFT and GOOG compute are cost centres; NVDA is the supplier



SPCX segment operating margin (% of segment revenue), quarter ended 2026-06-30. One segment is positive; the launch business is the worst-performing, not the AI buildout. Read on the filed sign — the served layer returns the consolidated line as +143 against a filed (143), and 16 of 20 served SPX operating-income facts are sign-stripped. Source: [SPCX 10-Q p.30](#).

The definition excludes exactly the costs a valuation would need. Nameplate compute draw is IT load — it "does not include power we install and use for our supporting infrastructure such as cooling systems, power distribution losses, lighting, security systems, or facility-level overhead." It is a capacity number, and no disclosed \$/kW converts it into revenue or earnings. This means the AI segment is carried at INVESTED CAPITAL — a cost basis — and that is a constitutional bound rather than an analytical gap, the only basis the disclosure supports. What it costs the reader: there is no AI valuation in this thesis, only a bound.

THE DEEP CUT · A THREE-LEVEL DECOMPOSITION NOBODY HAD OPENED

SPCX files a three-level revenue decomposition, and every level closes exactly

Nine artifacts inherited from 003 recorded these cuts as unformable. They are filed — under `srt:ProductOrServiceAxis`, [SPCX 10-Q p.13](#).

Level	Composition, H1 2026	Sum	Filed total	Closes
Space	Launch Services 978 + Launch & Development 603	1,581	1,581	exact
Connectivity	Consumer 4,633 + Enterprise & Government 2,915	7,548	7,548	exact
AI	AI Solutions & Infrastructure 2,669 + Advertising 710	3,379	3,379	exact
Σ segments	Space 1,581 + Connectivity 7,548 + AI 3,379	12,508	12,508	exact
By nature	Services 11,667 + Products 841	12,508	12,508	exact

Three independent closures, and the same 12,508 reached by two routes. That is the derivation's own check: **if the deep cut were a taxonomy the thesis had imposed, it would not close to the filer's own totals.** The axis is therefore **confirmed as a partition by arithmetic** — which matters, because XBRL does not declare a dimension exhaustive.

Reading the filing caught two transcription errors the artifacts had absorbed rather than surfaced. Enterprise & Government is 2,915, not 2,914 — and the wrong figure had produced a \$1M "rounding residual" that a defaults artifact wrote a whole section explaining how to treat. And the AI row's third opex line is the filed **Restructuring charges**, not an unnamed "Other".

This means the two errors are one error. Both carried a figure forward from a draft instead of re-reading the filing, and **both then built a construct to accommodate it** — a rounding default, a vague label. **A default is the one construct a reader cannot falsify**, being a stated reading rather than a filed claim. Every closure is now **exact: a chain carrying a rounding allowance cannot detect an error, because the allowance absorbs exactly the error it would surface.**

The capital goes to the segment that loses the most

One table, on the same page as the segment income table, states the allocation without ambiguity. Of \$18,369M of quarterly capital expenditure, **86.2% goes to AI** — the segment running the largest operating loss.

Segment	Capex, Q2 2026	Share	Operating result
AI	\$15,828M	86.2%	\$(1,257)M
Connectivity	\$1,367M	7.4%	+\$1,656M
Space	\$1,174M	6.4%	\$(542)M
Total	\$18,369M	100%	\$(143)M

The build is funded externally, and by an order of magnitude. Operating cash flow of +\$3,466M covers just **0.1005×** the **\$(34,487)M** investing outflow — against a **1.0×** falsifier. Financing supplied **+\$100,291M: \$85,675M** of IPO equity and **\$40,869M** of notes at a **6.03%** effective rate ([SPCX 10-Q p.50](#)).

This means CAPITAL is the binding constraint, not demand and not technology. Every segment's growth is downstream of a funding decision — and the issuer's own capex narrative names **data centers before launch facilities**, which the numbers confirm at **86.2% to 13.8%**. **What it means for the downside:** at this coverage the equity is not funding growth from operations, so the price of the next dollar is set by the capital markets rather than by the business.


And at this moment the constraint binds through one disclosure, which is now resolved. The Cursor acquisition is **complete** — not pending. On **2026-08-14** the merger became effective, converting Cursor equity into **389,289,254** Class A shares plus **1,752,426** shares for vested RSUs, against an **implied equity value of \$60.0 billion**. A further **29,128,326** RSUs and **44,365,047** options were assumed — **464,535,053 shares of total potential dilution** ([8-K 2026-08-14, Item 2.01](#)). **What this means for the share count every downstream thesis uses:** the registered count is **13,176,000,000** at 2026-06-30 and the pro-forma count is **13,567,041,680**. This anchor prices on the **pro-forma** basis, and a thesis using the 10-Q count alone **understates the denominator by 391 million shares** for a transaction that has already closed.

THE REVERSE-DCF · WHAT THE PRICE REQUIRES, SOLVED BACKWARDS

The live price requires perpetual growth equal to the required return

A forward DCF is not admissible here — consolidated free cash flow is negative and fails the programme's ≥3-year positive-FCF limb. The reverse-DCF is admissible precisely because it forecasts nothing: it solves for the growth a price implies and reports it.

Step	Value
Earnings base (annualisation, not a forecast)	\$6,624M
Implied multiple	312.8×
Rate spread — Gordon: $P/E = 1/(r - g)$	$r - g = 0.32\%$
Implied g at $r = 15\%$ and 25% — the constitution's band	14.68% – 24.68%
Implied terminal share of EV, against F1's 50–70% cap	~100%



Modeled segment value ranges, USD billions, each regime at the multiple its anchor names. Tops of all three sum to \$307bn against a \$2,071.8bn live market capitalisation — 14.8%. MODELED and deliberately wide, because no admissible comparator exists for any segment. Source: [SPCX 10-Q p.30](#).

$g \approx r - 0.32\%$ is the whole result in one line. At every point in the constitution's own 15–25% band, the growth the price implies is essentially the required return itself; when g approaches r , terminal value approaches the entire present value, and almost none of the \$2.07 trillion pays for cash flows anyone can currently see. This breaches F1's own terminal-value cap, and the anchor reports three readings rather than resolving them — the price may be too high; the cap may be the wrong instrument at this stage of a capital-intensive build; or the attribution may be the artefact, since the 313× sets Space and AI to zero by construction. What this means for the position: it produces neither a long nor a short thesis, but the demand the price makes on the business — that Connectivity compound at the cost of capital, indefinitely, with the other two segments contributing nothing.

The constitution's own anchor does not reconcile with the IPO price

The anchor price basis was checked rather than inherited, and it does not close. This is reported as a finding, because a downstream thesis inheriting the ~\$1.62T figure would inherit an unexplained 9% gap.

Basis	Value	Source or derivation
Shares outstanding, 2026-06-30	13,176,000,000	Class A 7,607M + Class B 5,569M — closes exactly
+ Cursor shares issued 2026-08-14	391,041,680	389,289,254 + 1,752,426 vested RSUs
Pro-forma shares	13,567,041,680	
Live price , close 2026-09-18	\$152.71	keyless NASDAQ quote, data_class fast
⇒ Live market capitalisation	≈ \$2.072T	13.567bn × \$152.71
IPO price, June 2026	\$135.00	638,888,888 shares; net proceeds \$85,675M (p.52)
Constitution anchor	~\$1.62T	struck 2026-06, basis not stated in this workspace

The unreconciled number is the anchor's implied price per share. $\$1.62\text{T} \div 13,176,000,000 = \122.95 , and on the pro-forma count $\$1.62\text{T} \div 13,567,041,680 = \119.41 . **The IPO priced at \$135.00.** So the constitution's anchor and the IPO price **differ by ~9% on the filed share count before any market move**, and the anchor's own basis is not recoverable from this workspace.

This means an earlier draft's headline spread was the wrong comparison. It reported the live price against the **IPO price** — $152.71/135.00 - 1 = +13.1\%$ — which measures the aftermarket, not the anchor. **Against the anchor's own implied price the live market is +27.9% on the pro-forma count** (\$152.71 vs \$119.41). Same two prices, two spreads, and only one of them answers the question the anchor poses.

What this thesis does with it. It publishes **both bases and resolves neither**, and records the anchor's reconciliation as UNRESOLVABLE-FROM-PUBLIC-SOURCES — **the remedy is the basis the constitution struck it on, which is a document input rather than a filing. What it means for downstream theses:** any thesis citing "the anchor" must state **which** anchor — the constitution's ~\$1.62T, the IPO's \$135.00, or the live \$152.71 — because they are three different numbers and the programme currently quotes all three.

WHAT COULD NOT BE VALUED · BOUNDS RECORDED, NEVER IMPLIED

Four bounds stand, and a fifth was retracted because the filing answers it

A bound is a result, not a gap. Each of the following was tested against the filing and could not be resolved from public sources; each is recorded with its disposition class, because **the classes have different remedies** — a research queue, a schema change, or nothing at all.

Item	Disposition class	Why it cannot be drawn
Aviation and maritime splits	UNRESOLVABLE-FROM-PUBLIC-SOURCES	they sit inside EnterpriseAndGovernmentMember, which the issuer files as one figure
Grok versus compute economics	UNRESOLVABLE-FROM-PUBLIC-SOURCES	both sit inside AI Solutions & Infrastructure; advertising alone is separable , at 21.0%
Subscribers and ARPU by channel	UNRESOLVABLE-FROM-PUBLIC-SOURCES	revenue decomposes into two channels; the denominator does not — ARPU is service-line only
The constitution's ~\$1.62T basis	UNRESOLVABLE-FROM-PUBLIC-SOURCES	it implies \$122.95 /share against an IPO at \$135.00 ; the basis is a document, not a filing
Segment-level capex	RETRACTED — it is filed	\$1,174 / \$1,367 / \$15,828 on the same page as the segment income table

The retraction is the substantive entry, and it is the most instructive thing in this report. An earlier version of this thesis's own synthesis recorded segment-level capex as UNRESOLVABLE-FROM-PUBLIC-SOURCES — *"company-level attribution only."* **That was false.** The figure is filed in Note 18's supplemental block, **on the same page as the segment income table the entire anchor is built on.**

This means the error was not a search failure — it was an artifact that did not open the page it was citing. And a bound asserted as unresolvable is the most expensive error available here: **it closes a question the filings answer, and no downstream thesis will re-open it.** The cost is not one missing row — **it is that the retracted row was load-bearing**: the 86.2% capex concentration is the anchor's capital-allocation finding, and it was unavailable for exactly as long as the bound stood.

What this means for how the programme should read a bound. Every remaining UNRESOLVABLE above was tested by **opening the page**, and each names the disclosure that would resolve it. **An unresolvable class without a named remedy is not a bound — it is an unexamined claim**, and this thesis treats the difference as the deliverable.

THE PILLARS · SIX VERDICTS REPORTED WITH THEIR FALSIFIER OUTCOMES

All six pillars hold, and three of them hold only as bounds

The thesis is specified as six pillars, each with a registered falsifier. They are reported here as verdicts rather than conclusions, because what a pillar does when its falsifier does not fire is different from what it establishes as a value.

Pillar	Verdict	What it turns on
PIL-1 Decomposition	HOLDS	3 of 3 segments file both a revenue line and an operating result
PIL-2 Operating leverage	HOLDS BOUND	+79.4% income on +65.8% revenue; DA-10 travels
PIL-3 AI valuation	BOUND ONLY	not a value — invested capital; advertising the only separable 21.0%
PIL-4 Launch value	HOLDS	ex-R&D result positive at +\$534M; numerator MODELED
PIL-5 Capital constraint	HOLDS	coverage 0.1005×, an order of magnitude below the falsifier
PIL-6 Comparability	HOLDS	partition admits 0 of 11; rows name 005, 006 and 009

The badges are not decoration — they distinguish three kinds of "holds". Five pillars hold as findings, one as a bound, and one as a finding carrying a bound. This means a downstream thesis reading "PIL-2 HOLDS" and pricing the leverage case without the DA-10 bound has read the verdict and missed the qualifier — exactly the inheritance failure the carry-the-basis rule exists to prevent.

The capability timeline is partial, and it is stated rather than implied. This thesis is market_data_stage: per_row — sotp-valuation and reverse-dcf run at late, the other five executed rows at none. It is the first thesis in the workspace with a live market-data dependency, and so the first where PRICE_ACCESS_PREMATURE could fire. Its nearest timeline equivalent is the seven-trigger register, which is datability-driven rather than capability-adoption-driven.

ADVERSARIAL REVIEW · RUN AGAINST THIS THESIS

Twelve findings were run against this thesis, four of them high

An adversarial pass was run against the finished thesis, and its findings are reported here rather than filed elsewhere — two of them **block the thesis's entry into the programme knowledge base**.

#	Severity	Finding
C-1	HIGH	The entity_claims schema names the field ticker ; every consumer reads entity . The index reports contradictions: 7 with entity='None' for all seven — the detector cannot see what it detects
C-2	HIGH	No retrieved_at field. A figure's observation date is recorded; its retrieval date is not — a stale datum cannot be told from a fresh one after the fact
C-3	HIGH	The 313× headline is attribution-dependent by construction — addressed by the allocation ladder, which the challenge pass did not have
C-4	HIGH	The AI segment is carried at invested capital with no AI valuation at all ; a reader may read a bound as a refusal to analyse
Eight further findings at MEDIUM — including the two transcription errors this report carries on the decomposition page, which were found by re-reading the filing rather than by any gate		

C-1 and C-2 are the same defect seen from two sides. The schema **names** a field one way and **omits** another entirely; both mean a consumer gets a value that **does not answer the question the field name asks**. **This means** they are one discipline, not two bugs: **a provenance field is only useful if the name, the value and the consumer agree about what it holds**.

What it means that this is reported at all. The challenge pass found defects that **twelve completed artifacts, three automated gates and a re-pack did not** — by reading for the consumer's expectation rather than the author's intent. **Both HIGH findings block knowledge-base entry**, and the anchor publishes with them open: **a downstream thesis that inherits these artifacts inherits the defect, and the only way it can know that is if the anchor says so**.

What 005, 006 and 009 may price off, and what they may not

The anchor exists so that four downstream theses do not each re-derive the same boundary. The permissions are explicit, and the prohibitions matter more than the permissions.

Downstream	May price off	May not
005	All three segment regimes; the 313× anchor fact; the allocation ladder	treat 313× as a valuation conclusion — it is a fact about where the price sits
006	The Connectivity row — SPCX is in 006's universe and appears in no other thesis's SPCX valuation	price off IRDM, GSAT, ASTS, VSAT or SATS — five excluded names, four of them in 006's own universe
009	The AI row as an invested-capital bound, with DA-11	construct an AI multiple from MSFT, GOOG, NVDA or VRT — cost centres and a supplier

The prohibitions are the load-bearing part, and the reason is that this failure is invisible from inside. A downstream thesis that prices SPCX's Connectivity off IRDM will look **internally consistent** — its arithmetic will reproduce, its sources will resolve, and every gate it runs will pass. **What it will have done is re-open a class this anchor closed**, and nothing in that thesis's own structure can detect it. **That is why the partition is published rather than the reasoning behind it alone.**

Seven triggers carry the anchor forward, each with a datum and a threshold rather than a direction — so a downstream thesis can check whether the anchor has gone stale without re-running it. **The highest-value resolution is identified:** the constitution's ~\$1.62T basis is unrecoverable from filings, and **one document input closes it.**

What this means for the programme. 001 established the quantities, 002 established that a grade does not carry a basis, 003 mapped where the margin sits, and **004 establishes the floor: what SPCX's parts are worth, and which of those parts cannot be valued at all.** Every one of the four rests on a refusal as much as a finding — and the anchor's refusal is that **the sector's benchmark company has no admissible comparators, so its value must be bounded rather than measured.**

The recommendation is not a position. This thesis produces no trades by design, and its deliverable is a citation rule: **carry the basis with the grade, and carry the boundary with the row.**

DISCLAIMER

This document is **research and analysis**, produced by an automated research system for internal use. It is **not investment advice**, and it is **not an offer or solicitation** to buy or sell any security.

It is based on the sources cited inline. Those sources may be incomplete, delayed, or wrong, and the analysis may have misread them; no representation or warranty is made as to accuracy or completeness. Figures are as of their stated as_of date and are not updated.

This document contains hypotheses, not established conclusions. Where a claim is a projection, an inference, or a judgement rather than a reported fact, it is marked as such. Statements about the future are forward-looking and inherently uncertain. Past performance is not indicative of future results.

Recipients must do their own due diligence and consult their own advisers before acting on anything here. The authors and distributors accept no liability for any loss arising from reliance on this document.

© agentii-space-tech-SPCX · 2026-09-18 · generated 2026-09-19 14:00 UTC